



中山大學
SUN YAT-SEN UNIVERSITY

Lecture 11

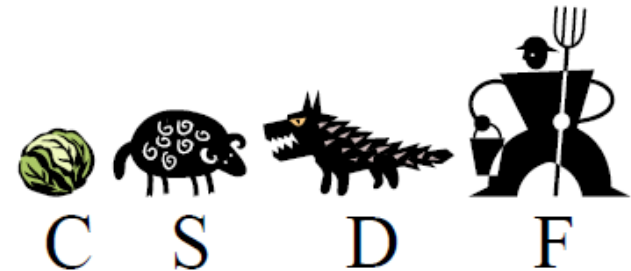
Search

Algorithm Design and Analysis

zhangzizhen@gmail.com

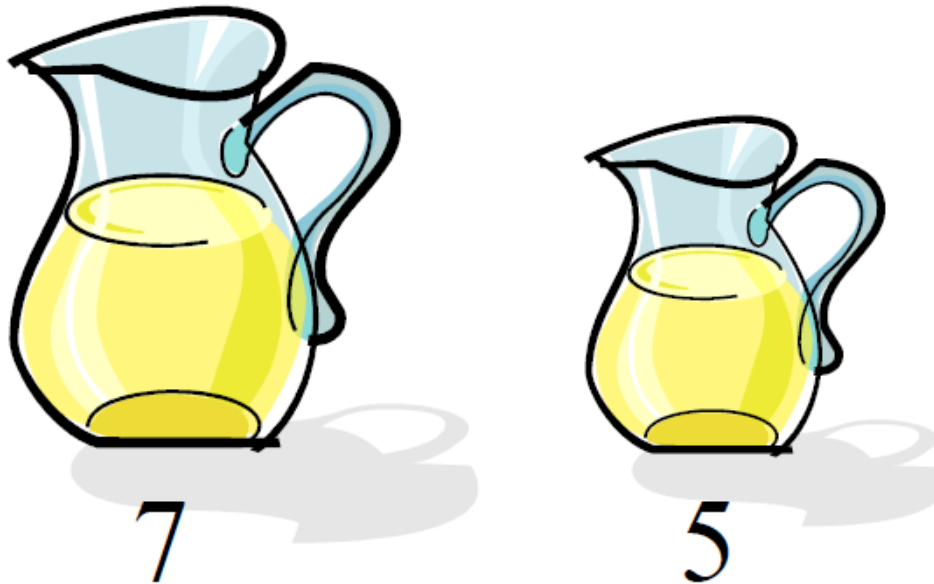
The search example

- State space \mathbf{S} : all valid configurations
- Initial states (nodes) $\mathbf{I} = \{(\mathbf{CSDF},)\} \subseteq \mathbf{S}$
 - Where's the boat?
- Goal states $\mathbf{G} = \{(\mathbf{,CSDF})\} \subseteq \mathbf{S}$
- Successor function $\mathbf{succs(s)} \subseteq \mathbf{S}$: states reachable in one step (one arc) from \mathbf{s}
 - $\mathbf{succs}((\mathbf{CSDF},)) = \{(\mathbf{CD}, \mathbf{SF})\}$
 - $\mathbf{succs}((\mathbf{CDF},\mathbf{S})) = \{(\mathbf{CD},\mathbf{FS}), (\mathbf{D},\mathbf{CFS}), (\mathbf{C}, \mathbf{DFS})\}$
- $\mathbf{cost(s,s')} = 1$ for all arcs. (weighted later)
- The search problem: find a solution path from a state in \mathbf{I} to a state in \mathbf{G} .
 - Optionally minimize the cost of the solution.



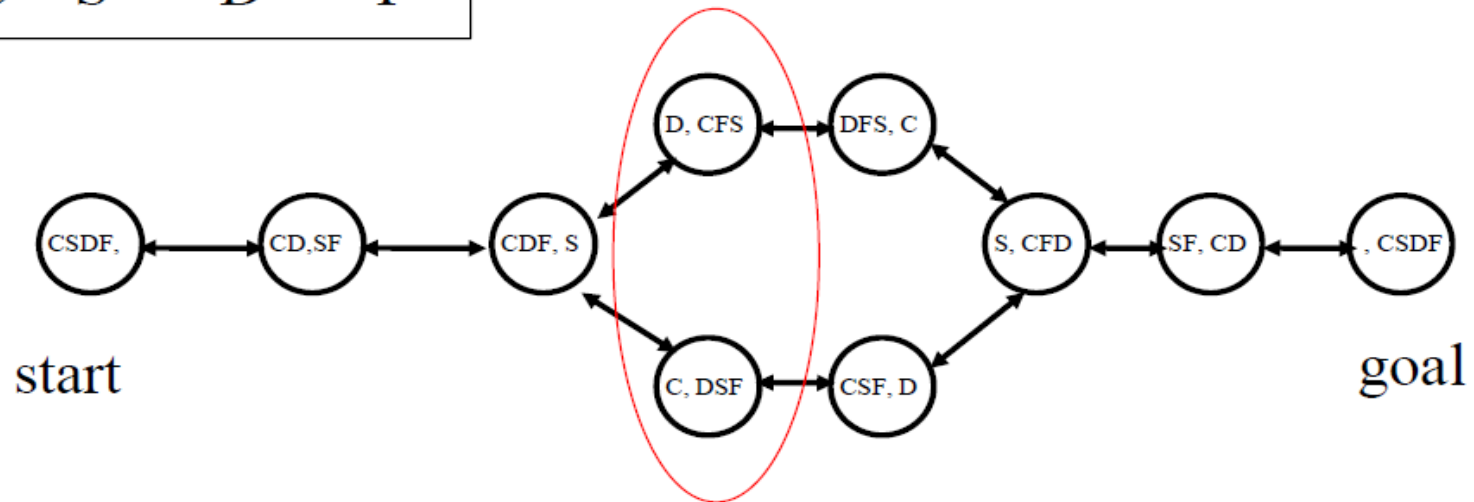
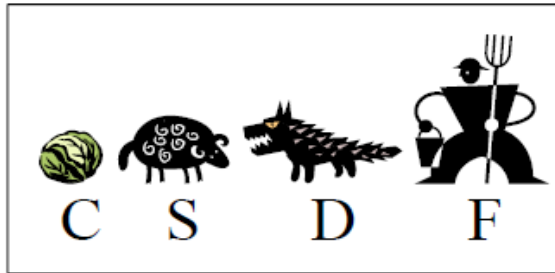
The search example

- Water jugs: how to get 1?



- Goal? (How many goal states?)
- Successor function: fill up (from tap or other jug), empty (to ground or other jug)

A directed graph in state space



- In general there will be many generated, but un-expanded states at any given time
- One has to choose which one to expand next
- Deep or shallow?

Uninformed search

- Uninformed means we only know:
 - The goal test
 - The **succs()** function
- But not which non-goal states are better: that would be informed search.
- For now, we also assume **succs()** graph is a **tree**.
 - Won't encounter repeated states.
 - We will discuss it later.
- Search strategies: BFS, UCS, DFS, IDS.

Breadth-first search (BFS)

- Use a queue (First-in First-out)

en_queue(Initial states)

While (queue not empty)

 s = de_queue()

 if (s==goal) success!


 T = succs(s)

 for t in T: t.prev=s

 en_queue(T)

endWhile

We need back
pointers to recover
the solution path.



Performance of BFS

- Assume:
 - the graph may be infinite.
 - Goal(s) exists and is only finite steps away.
- Will BFS find at least one goal?
- Will BFS find the least cost goal?
- Time complexity?
 - Number of states generated
 - Goal: d edges away
 - Branching factor: b
- Space complexity?
 - Number of states stored

Performance of BFS

- Completeness: yes, BFS will find a goal.
- Optimality: yes, if edges cost 1 (more generally positive non-decreasing in depth), no otherwise.
- Time complexity (worst case): goal is the last node at radius d .
 - Have to generate all nodes at radius d .
 - $b + b^2 + \dots + b^d \sim O(b^d)$
- Space complexity: $O(b^d)$

Uniform-cost search

- Find the least-cost goal
- Each node has a path cost from start (= sum of edge costs along the path). Expand the least cost node first.
- Use a priority queue instead of a normal queue
 - Always take out the least cost item
 - Remember *heap*? time $O(\log(\text{number of items in heap}))$
- Complete and optimal (if edge costs $\geq \varepsilon > 0$)
- Time and space: can be much worse than BFS
 - Let C^* be the cost of the least-cost goal
 - $O(b^{C^*/\varepsilon})$, possibly $C^*/\varepsilon \gg d$


Depth-first search

- Use a stack (First-in Last-out)

```

push(Initial states)
While (stack not empty)
    s = pop()
    if (s==goal) success!
    T = succs(s)
    push(T)
endWhile
  
```

This is non-recursive
 implementation of DFS, recursive
 implementation is more common




Performance of DFS

- m = maximum depth of graph from start
- Space complexity: $O(mb)$
- Infinite tree: may not find goal (incomplete)
- May not be optimal
- Finite tree: may visit almost all nodes, time complexity $O(b^m)$

Iterative deepening search

1. DFS, but stop if path length > 1 .
2. If goal not found, repeat DFS, stop if path length > 2 .
3. And so on...

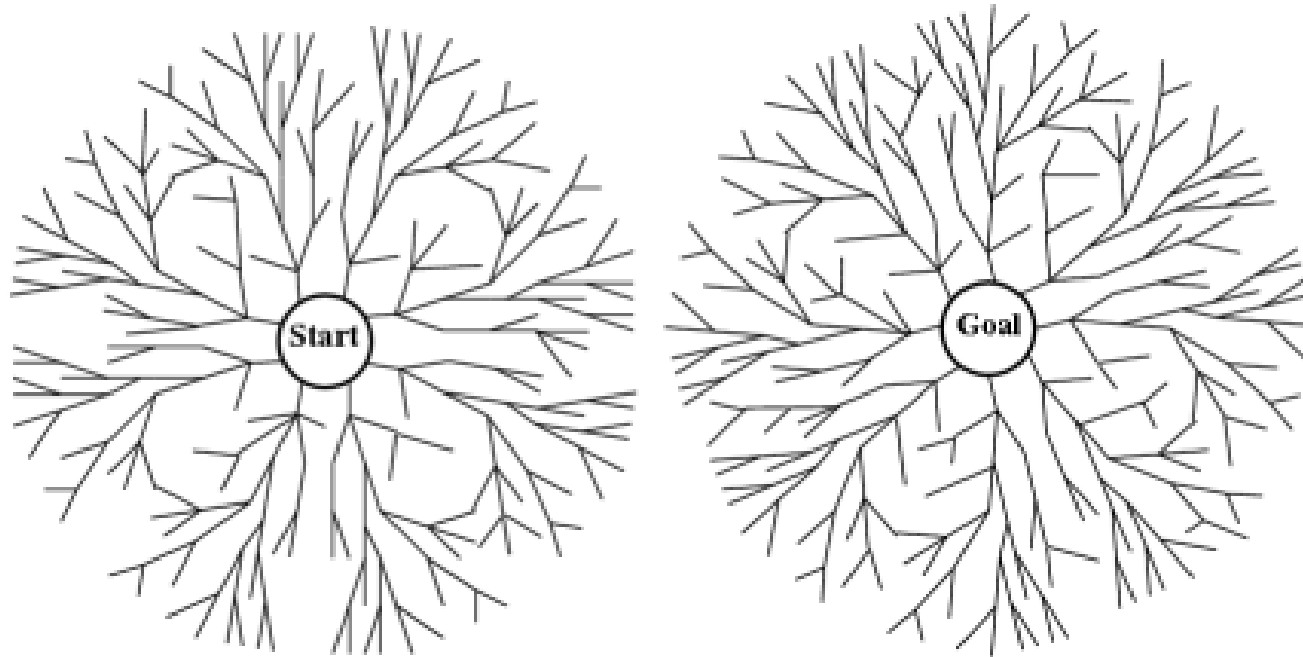


Iterative deepening search

- BFS + DFS
 - Complete, optimal like BFS
 - Small space complexity like DFS
- •A huge waste?
 - Each deepening repeats DFS from the beginning
 - No! $db + (d-1)b^2 + (d-2)b^3 + \dots + b^d \sim O(b^d)$
 - Time complexity like BFS

Bidirectional search

- Breadth-first search from both start and goal
- Stop when fringes meet
- The fringes(边缘) are $O(b^{d/2})$
- Generates $O(b^{d/2})$ instead of $O(b^d)$ nodes



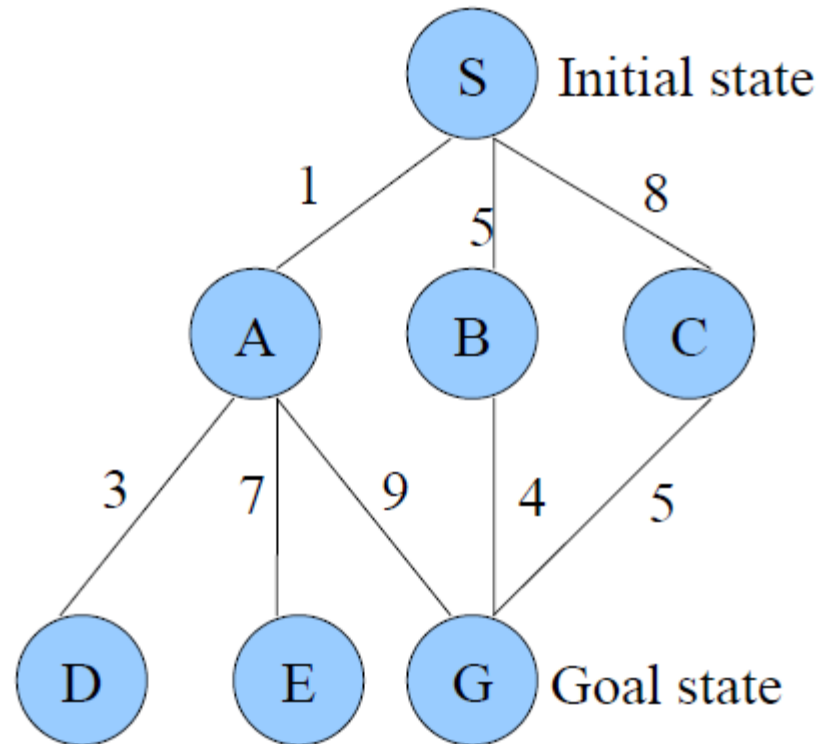
Performance of search algorithms

b: branching factor (assume finite) d: goal depth m: graph depth

	Complete	optimal	time	space
Breadth-first search	Y	Y, if ¹	$O(b^d)$	$O(b^d)$
Uniform-cost search	Y	Y	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$
Depth-first search	N	N	$O(b^m)$	$O(bm)$
Iterative deepening	Y	Y, if ¹	$O(b^d)$	$O(bd)$
Bidirectional search	Y	Y, if ¹	$O(b^{d/2})$	$O(b^{d/2})$

1. edge cost constant, or positive non-decreasing in depth

Search example



- All edges are directed, pointing downwards

Node expansion

- Depth-First Search:

- S A D E G
- Solution found: S A G

- Breadth-First Search:

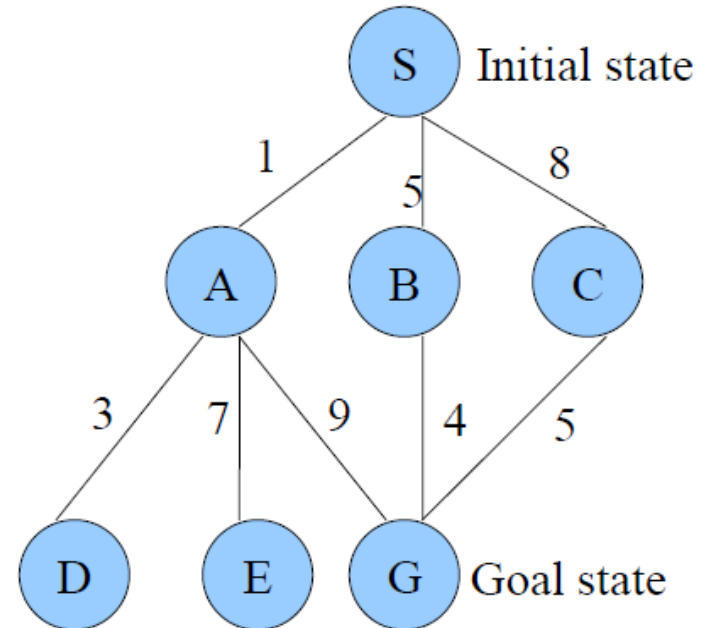
- S A B C D E G
- Solution found: S A G

- Uniform-Cost Search:

- S A D B C E G
- Solution found: S B G (This is the only uninformed search that worries about costs.)

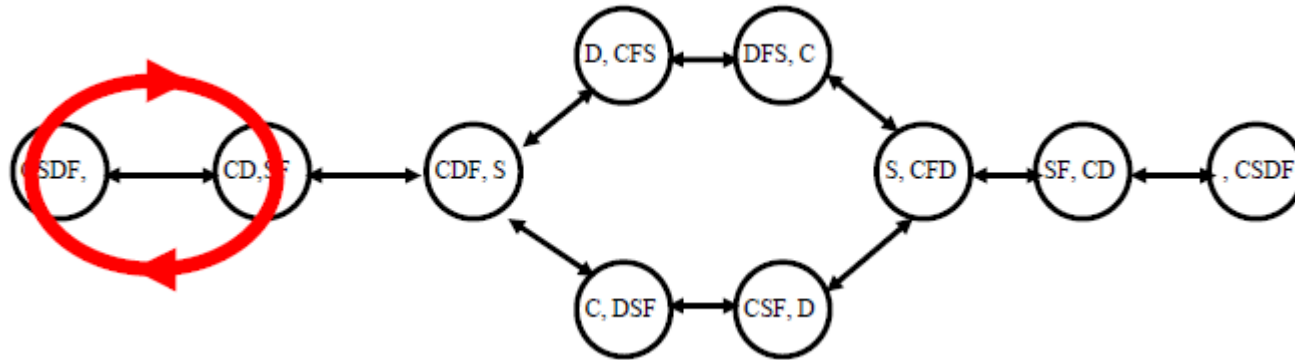
- Iterative-Deepening Search:

- S A B C S A D E G
- Solution found: S A G



General graph search

- The problem: repeated states



- We have to remember already-expanded states (**CLOSED**).
- When we take out a state from the fringe (**OPEN**), check whether it is in **CLOSED** (already expanded).
 - If yes, throw it away.
 - If no, expand it (add successors to **OPEN**), and move it to **CLOSED**.

General graph search

- BFS:
 - Still $O(b^d)$ space complexity
- DFS:
 - Memorizing DFS (MEMDFS): memorize every expanded states
 - Path Check DFS (PCDFS): remember only expanded states on current path (from start to the current node)

Informed search

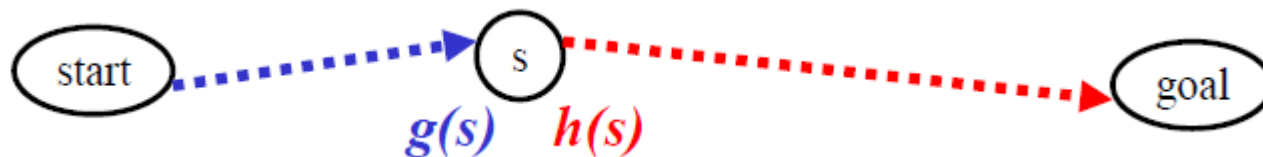
- Uninformed search

- Knows the actual path cost $g(s)$ from the start to a node s .



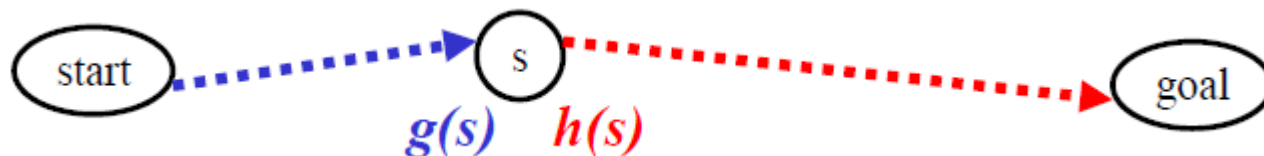
- Informed search

- Also has a heuristic $h(s)$ of the cost from s to goal.
- Can be much faster than uninformed search.



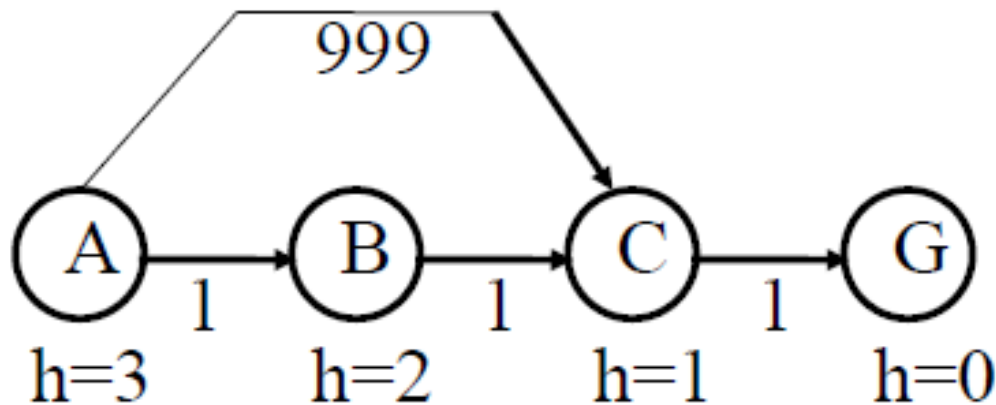
Recall: Uniform-cost search

- Uniform-cost search: uninformed search when edge costs are not the same.
- Complete (will find a goal). Optimal (will find the least-cost goal).
- Always expand the node with the least $g(s)$
- Use a priority queue:
 - Push in states with their first-half-cost $g(s)$
 - Pop out the state with the least $g(s)$ first
- Now we have an estimate of the second-half-cost $h(s)$, how to use it?



Best-first greedy search


- Use $h(s)$ instead of $g(s)$
- Always expand the node with the least $h(s)$
- Not optimal



It will follow the path A \rightarrow C \rightarrow G


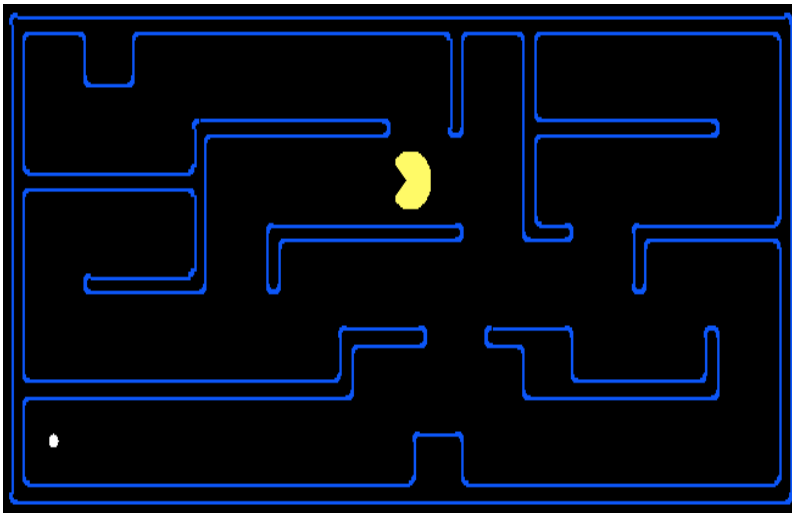
A search

- Use $f(s)=g(s)+h(s)$
- Always expand the node with the least $g(s)+h(s)$
- A search is not always optimal



A* search

- Same as A search, but the heuristic function $h()$ has to satisfy $h(s) \leq h^*(s)$, where $h^*(s)$ is the true cost from node s to the goal.
- Such heuristic function $h()$ is called **admissible**.
- An admissible heuristic never over-estimates
- A search with admissible $h()$ is called **A* search**.



Admissible heuristic functions h

- 8-puzzle example

Example State

1		5
2	6	3
7	4	8

Goal State

1	2	3
4	5	6
7	8	

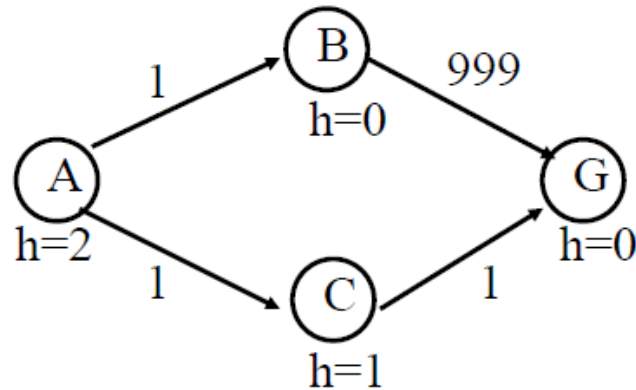
- Which of the following are admissible heuristics?
 - $h(n)$ =number of tiles in wrong position
 - $h(n)=0$
 - $h(n)=1$
 - $h(n)$ =sum of Manhattan distance between each tile and its goal location

Admissible heuristics


- A heuristic function h_2 **dominates** h_1 if for all s
 $h_1(s) \leq h_2(s) \leq h^*(s)$
- $d = 14$,
 - $A^*(h_1) = 539$ nodes
 - $A^*(h_2) = 113$ nodes
- $d = 24$,
 - $A^*(h_1) = 39,135$ nodes
 - $A^*(h_2) = 1,641$ nodes
- We prefer heuristic functions as close to h^* as possible, but not over h^* .
- Good heuristic function might need complex computation
- Time may be better spent, if we use a faster, simpler heuristic function and expand more nodes.

Some tricks

- A* should terminate only when a goal is popped from the priority queue



- A* can revisit an expanded state, and discover a shorter path



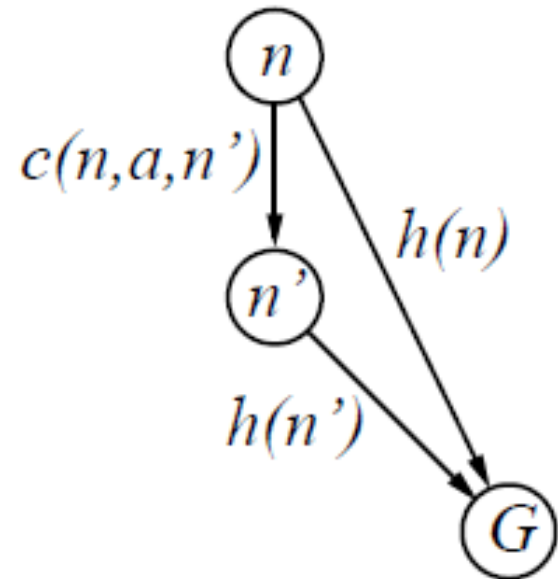
The A* algorithm

1. Put the start node **S** on the priority queue, called **OPEN**
2. If **OPEN** is empty, exit with failure
3. Remove from **OPEN** and place on **CLOSED** a node **n** for which **f(n)** is minimum
4. If **n** is a goal node, exit (trace back pointers from **n** to **S**)
5. Expand **n**, generating all its successors and attach to them pointers back to **n**. For each successor **n'** of **n** not on **CLOSED**
 1. If **n'** is not already on **OPEN**, estimate $h(n')$, $g(n')=g(n)+c(n,n')$, $f(n')=g(n')+h(n')$, and place it on **OPEN**.
 2. If **n'** is already on **OPEN**, then check if $g(n')$ is lower for the new version of **n'**. If so, then:
 - Redirect pointers backward from **n'** along path yielding lower $g(n')$.
 - Put **n'** on **OPEN**.
 - If $g(n')$ is not lower for the new version, do nothing.
6. Goto 2.

Consistent heuristics


- Consistency is analogous to the triangle inequality from Euclidian geometry.
- A heuristic is **consistent** if $h(n) \leq c(n, a, n') + h(n')$
- If h is consistent, then for every child n' of n , we have:

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n, a, n') + h(n') \\ &\geq g(n) + h(n) \\ &= f(n) \end{aligned}$$
- That is, $f(n)$ is non-decreasing along any path.



Behavior of A^* with consistent heuristic

- If h is consistent, then A^* expands nodes in order of increasing f value. In such a case, A^* can be implemented more efficiently — no node needs to be processed more than once.
- Gradually adds “ f -contours” of nodes (breadth-first adds layers)
- Contour i has all nodes with $f = f_i$, where $f_i < f_{i+1}$

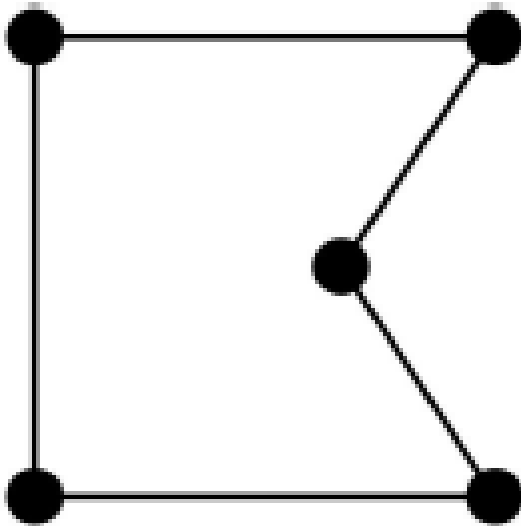


Properties of A*

- Complete? Yes
- Time? $O(\text{entire state space})$ in worst case, $O(d)$ in best case
- Space? Keeps all nodes in memory
- Optimal? Yes

Traveling salesman problem

- For a node s
 $f(s)=g(s)+h(s)$
- What is $g(s)$, $h(s)$?
- A* v.s. branch-and-bound



Iterative-Deepening A*

function IDA*(*problem*) returns a solution

inputs: *problem*, a problem

$f_0 \leftarrow h(\text{initial state})$

for $i \leftarrow 0$ to ∞ do

$\text{result} \leftarrow \text{COST-LIMITED-SEARCH}(\text{problem}, f_i)$

 if result is a solution then return result

 else $f_{i+1} \leftarrow \text{result}$

end

function COST-LIMITED-SEARCH(*problem*, f_{\max}) returns solution or number

depth-first search, backtracking at every node n such that $f(n) > f_{\max}$

if the search finds a solution then

 return the solution

else

 return $\min\{f(n) \mid \text{the search backtracked at } n\}$

Thank you!

